

Summary for Introduction to Machine Learning 2019

Regression: Predict real valued labels

Linear Regression

$f(x) = w_1x_1 + \dots + w_dx_d + w_0 = \tilde{w}^T \tilde{x}$ with $\tilde{w} = [w_1 \dots w_d, w_0]$ and $\tilde{x} = [x_1 \dots x_d, 1]$
Residual: $r_i = y_i - w^T x_i$, $x_i \in \mathbb{R}^d$, $y_i \in \mathbb{R}$
Cost / Objective function (is convex):
 $\hat{R}(w) = \sum_{i=1}^n r_i^2 = \sum_{i=1}^n (y_i - w^T x_i)^2$
Optimal weights:
 $w^* = \underset{w}{\operatorname{argmin}} \sum_{i=1}^n (y_i - w^T x_i)^2$

Closed form solution: $w^* = (X^T X)^{-1} X^T y$
Gradient: $\nabla_w \hat{R}(w) = [\frac{\delta}{\delta w_1} \hat{R}(w) \dots \frac{\delta}{\delta w_d} \hat{R}(w)] = -2 \sum_{i=1}^n r_i x_i^T$
Non-linear functions: $f(x) = \sum_{i=1}^D w_i \phi_i(x)$

Convex function

$f: \mathbb{R}^d \rightarrow \mathbb{R}$ is convex $\Leftrightarrow x_1, x_2 \in \mathbb{R}^d, \lambda \in [0, 1]:$
 $f(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda f(x_1) + (1 - \lambda)f(x_2)$

Gradient Descent

1. Start at an arbitrary $w_0 \in \mathbb{R}^d$
2. For $t = 1, 2, \dots$ do $w_{t+1} = w_t - \eta_t \nabla \hat{R}(w_t)$

Gaussian/Normal Distribution

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right)$$

Multivariate Gaussian

$$f(x) = \frac{1}{2\pi\sqrt{|\Sigma|}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)}$$
$$\Sigma = \begin{pmatrix} \sigma_1^2 & \sigma_{12} \\ \sigma_{21} & \sigma_2^2 \end{pmatrix}, \quad \mu = \begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix}$$

Empirical risk minimization

Assumption: Data set generated iid from unknown distribution $P: (x_i, y_i) \sim P(X, Y)$.
True risk: $R(w) = \int P(x, y)(y - w^T x)^2 dx dy = \mathbb{E}_{x,y}[(y - w^T x)^2]$
Empirical risk:
 $\hat{R}_D(w) = \frac{1}{|D|} \sum_{(x,y) \in D} (y - w^T x)^2$
Generalization error: $|R(w) - \hat{R}_D(w)|$
Uniform convergence:
 $\sup_w |R(w) - \hat{R}_D(w)| \rightarrow 0$ as $|D| \rightarrow 0$

In general, it holds that:
 $\mathbb{E}_D[\hat{R}_D(\hat{w}_D)] \leq \mathbb{E}_D[R(\hat{w}_D)]$, where
 $\hat{w}_D = \underset{w}{\operatorname{argmin}} \hat{R}_D(w)$.

Cross-validation

For each model m

For $i = 1:k$

1. Split data: $D = D_{train}^{(i)} \uplus D_{val}^{(i)}$
2. Train model: $\hat{w}_{i,m} = \underset{w}{\operatorname{argmin}} \hat{R}_{train}^{(i)}(w)$

3. Estimate error: $\hat{R}_m^{(i)} = \hat{R}_{val}^{(i)}(\hat{w}_{i,m})$

After all iterations, select model:

$$\hat{m} = \underset{m}{\operatorname{argmin}} \frac{1}{k} \sum_{i=1}^k \hat{R}_m^{(i)}$$

Ridge regression

Regularization(corresponds to MAP estimation):
 $\min_w \frac{1}{n} \sum_{i=1}^n (y_i - w^T x_i)^2 + \lambda \|w\|_2^2 =$
 $\underset{w}{\operatorname{argmax}} P(w) \Pi_i P(y_i | x_i w)$

Sparse regression (L1, convex) encourages coefficients to be exactly 0 - automatic feature

selection

Closed form solution: $\hat{w} = (X^T X + \lambda I)^{-1} X^T y$
Gradient: $\nabla_w (\frac{1}{n} \sum_{i=1}^n (y_i - w^T x_i)^2 + \lambda \|w\|_2^2) = \nabla_w \hat{R}(w) + 2\lambda w$

Standardization

Goal: each feature: $\mu = 0, \sigma^2 = 1$:

$$\tilde{x}_{i,j} = \frac{(x_{i,j} - \hat{\mu}_j)}{\hat{\sigma}_j}$$

$$\hat{\mu}_j = \frac{1}{n} \sum_{i=1}^n x_{i,j}, \quad \hat{\sigma}_j^2 = \frac{1}{n} \sum_{i=1}^n (x_{i,j} - \hat{\mu}_j)^2$$

Dimension Reduction in unsupervised learning

Principal Component Analysis (linear)

Given $D \subseteq \mathbb{R}^d, 1 \leq k \leq d, \Sigma = \frac{1}{n} \sum_{i=1}^n x_i x_i^T, \mu = \frac{1}{n} \sum_i x_i = 0$ (data is centered)
 $(W, z_1, \dots, z_n) = \underset{W}{\operatorname{argmin}} \sum_{i=1}^n \|W z_i - x_i\|_2^2$
where $W \in \mathbb{R}^{d \times k}$ is orthogonal, $z_1, \dots, z_n \in \mathbb{R}^k$ is given by $W = (v_1 | \dots | v_k)$ and $z_i = W^T x_i = f(x)$ where $\Sigma = \sum_{i=1}^d \lambda_i v_i v_i^T$ where $\lambda_1 \geq \dots \geq \lambda_d \geq 0$
The projection is chosen to minimize the reconstruction error, choose k such that most of the variance is explained (like k-means)

Kernel PCA (nonlinear)

For $k = 1$: Kernel PCA

$$\alpha^* = \underset{\alpha^T K \alpha = 1}{\operatorname{argmax}} \alpha^T K^T K \alpha$$

With $K = \sum_{i=1}^n \lambda_i v_i v_i^T$ ($\lambda_1 \geq \dots \geq \lambda_d \geq 0$)

$$\alpha^* = \frac{1}{\sqrt{\lambda_1}} v_1$$

For general k : Kernel PCA

The kernel principal components are given by $\alpha^{(1)}, \dots, \alpha^{(k)} \in \mathbb{R}^n$
 $\alpha^{(i)} = \frac{1}{\sqrt{\lambda_i}}$ with $K = \sum_{i=1}^n \lambda_i v_i v_i^T$

A new point x is projected as z ,

$$z_i = \sum_{j=1}^n \alpha_j^{(i)} k(x, x_j)$$

Kernel-PCA corresponds to applying PCA in the feature space induced by the kernel k .

centering a kernel: $K' = K - KE - EK + EKE$

where $E = \frac{1}{n} [1, \dots, 1][1, \dots, 1]^T$

- complexity grow with number of data points, requires data specified as kernel

Autoencoders

Goal: learn identity function $x \approx f(x; \theta)$

$$f(x; \theta) = f_{dec}(f_{enc}(x; \theta_1); \theta_2)$$

NN autoencoders are ANNs where one output unit for each of d input units, nr of hidden units smaller than nr of inputs. Optimize w s.t. output agrees with input.

If activation func. is the identity, fitting NN autoencoder is equivalent to PCA.

Decision Theory

Bayesian Decision Theory

Given: $P(y|x)$, set of actions A and cost function $C: Y \times A \rightarrow \mathbb{R}$

$$a^* = \underset{a \in A}{\operatorname{argmin}} \mathbb{E}_y [C(y, a) | x] \quad (\text{cost for prediction } y)$$

when true label is a)

for logistic

$$\text{regression: } \underset{y}{\operatorname{argmax}} P(y|x) = \operatorname{sign}(w^T x)$$