

1 General

P-Norm: $\|x\|_p = (\sum_{i=1}^n |x_i|^p)^{\frac{1}{p}}$

Frobenious Norm: $\|A\|_F = \sqrt{\sum_{i,j} a_{ij}^2}$

Derivation rules: Chain rule: $D(f(g(x))) = Df(g(x)) * Dg(x)$

positive definiteness: A is p.s.d., then A is a real symmetric matrix and $x^T A x \geq 0$ for all x

Joint distribution: X, Y are RVs $f_{X,Y}(x, y) = \mathbb{P}(X \leq x, Y \leq y)$

Joint density: $f_{X,Y}(x, y) = \frac{\delta^2 F}{\delta x \delta y}(x, y)$

Conditional Probability: $\mathbb{P}(A|B) = \frac{\mathbb{P}(A \cap B)}{\mathbb{P}(B)}$

Law of total probability: $\mathbb{P}(B) = \sum_{i=1}^n \mathbb{P}(B|A_i) \mathbb{P}(A_i)$

Bayes rule: $\mathbb{P}(A|B) = \mathbb{P}(B|A) \frac{\mathbb{P}(A)}{\mathbb{P}(B)}$

Variance: $Var(X) = \mathbb{E}[(X - \mathbb{E}[X])^2] = \mathbb{E}[X^2] - \mathbb{E}[X]^2 \geq 0$

Convexity: A twice differentiable function $f: \mathbb{R}^d \rightarrow \mathbb{R}$ is convex iff for any $x \in \mathbb{R}^d$ its Hessian is p.s.d

Convex functions are closed under addition

2 Regression: Predict real valued labels

Linear Regression: Goal: Measure distance between predicted and target values $f(x) = w_1 x_1 + \dots + w_d x_d + w_0 = \tilde{w}^T \tilde{x}$ with $\tilde{w} = [w_1 \dots w_d, w_0]$ and $\tilde{x} = [x_1 \dots x_d, 1]$

Residual: $r_i = y_i - w^T x_i, x_i \in \mathbb{R}^d, y_i \in \mathbb{R}$

Cost / Objective function (is convex): $\hat{R}(w) = \sum_{i=1}^n r_i^2 = \sum_{i=1}^n (y_i - w^T x_i)^2$

Optimal weights: $w^* = \underset{w}{\operatorname{argmin}} \sum_{i=1}^n (y_i - w^T x_i)^2$

Closed form solution: $w^* = (X^T X)^{-1} X^T y$

Gradient: $\nabla_w \hat{R}(w) = [\frac{\delta}{\delta w_1} \hat{R}(w) \dots \frac{\delta}{\delta w_d} \hat{R}(w)] = -2 \sum_{i=1}^n r_i x_i^T$

Non-linear functions: $f(x) = \sum_{i=1}^D w_i \phi_i(x)$

Fisher consistency: Given a surrogate loss function $\psi: Y \times S \rightarrow \mathbb{R}$, the surrogate is said to be consistent with respect to the loss $L: Y \times S \rightarrow \mathbb{R}$, if every minimizer f of the surrogate risk function $R_\psi(f)$ is also a minimizer of the risk function $R_L(f)$. E.g. the hinge and the logistic losses are consistent with respect to the 0-1 loss.

Classification losses: $L_{\text{perceptron}}: \{-1, 1\} \times \mathbb{R} \rightarrow \mathbb{R}: y, f(x) \rightarrow \max(0, -yf(x))$

Find the best separation hyperplane

$L_{\text{hinge}}: \{-1, 1\} \times \mathbb{R} \rightarrow \mathbb{R}: y, f(x) \rightarrow \max(0, 1 - yf(x))$

Find large separation margin

$L_{\text{perceptron}}: \{-1, 1\} \times \mathbb{R} \rightarrow \mathbb{R}: y, f(x) \rightarrow \max \log(1 + \exp(-yf(x)))$

Link to cross entropy and probabilistic interpretation

Classification: Accuracy: $\frac{TP+TN}{TP+TN+FP+FN}$

Recall/Sensitivity/True positive rate/TPR: $\frac{TP}{TP+FN}$

Specify or True negative rate/TNR: $\frac{TN}{TN+FP}$

F1 score: $2 * \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$

Convex function: $f: \mathbb{R}^d \rightarrow \mathbb{R}$ is convex $\Leftrightarrow x_1, x_2 \in \mathbb{R}^d, \lambda \in [0, 1]: f(\lambda x_1 + (1-\lambda)x_2) \leq \lambda f(x_1) + (1-\lambda)f(x_2)$

Gradient Descent: 1. Start at an arbitrary $w_0 \in \mathbb{R}^d$

2. For $t = 1, 2, \dots$ do $w_{t+1} = w_t - \eta_t \nabla \hat{R}(w_t)$

Gaussian/Normal Distribution: $f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp(-\frac{(x-\mu)^2}{2\sigma^2})$

Multivariate Gaussian: $f(x) = \frac{1}{2\pi\sqrt{|\Sigma|}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)}$

$\Sigma = \begin{pmatrix} \sigma_{11}^2 & \sigma_{12} \\ \sigma_{21} & \sigma_{22}^2 \end{pmatrix}, \mu = \begin{pmatrix} \mu_1 \\ \mu_2 \end{pmatrix}$

Empirical risk minimization: Assumption: Data set generated iid from unknown distribution $P: (x_i, y_i) \sim P(X, Y)$.

True risk: $R(w) = \int P(x, y)(y - w^T x)^2 dx dy = \mathbb{E}_{x,y}[(y - w^T x)^2]$

Empirical risk: $\hat{R}_D(w) = \frac{1}{|D|} \sum_{(x,y) \in D} (y - w^T x)^2$

Generalization error: $|R(w) - \hat{R}_D(w)|$

Uniform convergence: $\sup_w |R(w) - \hat{R}_D(w)| \rightarrow 0$ as $|D| \rightarrow 0$

In general, it holds that: $\mathbb{E}_D[\hat{R}_D(\hat{w}_D)] \leq \mathbb{E}_D[R(\hat{w}_D)]$, where $\hat{w}_D = \underset{w}{\operatorname{argmin}} \hat{R}_D(w)$.

Cross-validation: For each model m For $i = 1:k$

1. Split data: $D = D_{\text{train}} \uplus D_{\text{val}}$

2. Train model: $\hat{w}_{i,m} = \underset{w}{\operatorname{argmin}} \hat{R}_{\text{train}}^{(i)}(w)$

3. Estimate error: $\hat{R}_m^{(i)} = \hat{R}_{\text{val}}^{(i)}(\hat{w}_{i,m})$

After all iterations, select model: $\hat{m} = \underset{m}{\operatorname{argmin}} \frac{1}{k} \sum_{i=1}^k \hat{R}_m^{(i)}$

Ridge regression: Regularization (corresponds to MAP estimation): $\min_w \frac{1}{n} \sum_{i=1}^n (y_i - w^T x_i)^2 + \lambda \|w\|_2^2 = \underset{w}{\operatorname{argmax}}_w P(w) \Pi_i P(y_i | x_i w)$

Sparse regression (L1, convex) encourages coefficients to be exactly 0 - automatic feature selection

Closed form solution: $\hat{w} = (X^T X + \lambda I)^{-1} X^T y$

Gradient: $\nabla_w (\frac{1}{n} \sum_{i=1}^n (y_i - w^T x_i)^2 + \lambda \|w\|_2^2) = \nabla_w \hat{R}(w) + 2\lambda w$

Standardization: Goal: each feature: $\mu = 0$,

$\sigma^2 = 1: \tilde{x}_{i,j} = \frac{x_{i,j} - \hat{\mu}_j}{\hat{\sigma}_j}$

$\hat{\mu}_j = \frac{1}{n} \sum_{i=1}^n x_{i,j}, \hat{\sigma}_j^2 = \frac{1}{n} \sum_{i=1}^n (x_{i,j} - \hat{\mu}_j)^2$

3 Classification

$h(x) = \text{sign}(w^T x)$

Losses: 0/1 loss: $\ell_{0/1}(w; x, y) = [y \neq \text{sign}(w^T x)]$

Perceptron loss: $\ell_p(w; x, y) = \max(0, -yw^T x)$

Hinge loss: $\ell_H(w; x, y) = \max(0, 1 - yw^T x)$

$\nabla_w \ell_p(w, x_i, y_i) = \begin{cases} 0, & \text{if } w^T x_i y_i \geq 0 \\ -y_i x_i & \text{else} \end{cases}$

SGD: GD requires sum over all data, slow for large datasets.

1. Choose random initial $w_0 \in \mathbb{R}^d$

2. For $t = 1, 2, \dots$ do:

(a) Choose $(x, y) \in D$ u.a.r (w/ replacement)

(b) Set $w_{t+1} = w_t - \eta_t \nabla \ell(w_t; x, y)$

SGD converges if $\sum_t \eta_t = \infty$ and $\sum_t \eta_t^2 < \infty$.

Mini-batch: Choose multiple datapoints at random; may converge faster.

Perceptron: SGD with ℓ_p and $\eta = 1$.

$\hat{w} = \arg \min_w \frac{1}{n} \sum_{i=1}^n \ell_p(w; x_i, y_i)$

If data linearly separable finds separator.

SVM: SGD with ℓ_H and regularization.

$\hat{w} = \arg \min_w \frac{1}{n} \sum_{i=1}^n \ell_H(w; x_i, y_i) + \lambda \|w\|_2^2$

$w_{t+1} = w_t(1 - 2\eta_t \lambda) + \eta_t x_i y_i [y_i w^T x_i < 1]$

Often $\eta_t = \frac{1}{\lambda t}$. **Works on non-linearly separable data**, finds best separator w.r.t. ℓ_H .

4 ???orthogonal distance

: Let $w \in \mathbb{R}^d$ and $H = \{x \in \mathbb{R}^d | \langle w, z \rangle = 0\}$ be a hyperplane. The orthogonal distance of a point $z \in \mathbb{R}^d$ to H can be computed as $\frac{|\langle w, z \rangle|}{\|w\|}$. Specifically, if w is a unit vector, the inner product $\langle w, z \rangle$ directly gives the distance of z to H .

5 Feature selection

Naive: try all subsets, and pick best (via cross-validation)

Greedy: Greedily add (or remove) features to maximize cross-validated prediction accuracy w.r.t. cost $c: V \Rightarrow \mathbb{R}$ of using features in subset of V . Forward (start with empty set) is faster, but backward is more resilient to "dependent" features. Applies to **any method**, but **slow** b/c trains many models and can be **suboptimal**.

L1-Regularization: regularize loss with $\|w\|_1$, automatic feature selection. Can be used for regression (Lasso), classification (L1-SVM) by replacing $\|w\|_2^2$ with $\|w\|_1$. Only works for **linear** models, but is **fast**.

6 Class imbalance

Downsample **loses data** but **fast**, upsample **random perturbation maybe unsafe**. Use cost-sensitive metrics controlling tradeoff: $\ell_{CS} = c_y \ell(w; x, y)$

Pred. lab.	True label			
		Positive	Negative	Σ
	Positive	TP	FP	p_+
Negative	FN	TN	p_-	
Σ	n_+	n_-		

Accuracy **bad** metric for imbalanced data.

Accuracy $\frac{TP+TN}{TP+TN+FP+FN}$ Precision

TPR, Recall $\frac{TP}{TP+FN}$ F1 score

FPR $\frac{FP}{TN+FP}$

Precision Recall Curve: Precision (y-axis) vs. Recall (x-axis).

Precision = 1 and Recall = 1 is optimal. Area under curve (AUC) can be used for comparison of algos.

Receiver Operator Characteristic (ROC) Curve: TPR (y-axis) vs. FPR (x-axis).

Random guessing achieves TPR = FPR line. TPR > FPR is better than random guessing. TPR = 1 and FPR = 0 is optimal. Area under curve (AUC) can be used for comparison of algos.

Theorem: Alg 1 dominates Alg 2 in terms of ROC curve \Leftrightarrow Alg 1 dominates Alg 2 in terms of Precision Recall curve.

One-vs-all: c classifiers, one for each class, pick highest confidence. Class may **not be lin. sep.** from all others. Note **Scaling + Imbalance**.

One-vs-one $c(c-1)/2$ classifiers, voting scheme with highest number of positive prediction wins: **no confidence needed**.

Ideally $\mathcal{O}(\log c)$ classifiers, theoretical optimum.

Multi-class SVM: Maintains c weight vectors, want $w^{(y)T} x \geq \max\{w^{(i)T} x\} + 1$ for correct label y .

$\ell_{MC-H}(w^{(1)}, \dots, w^{(c)}; x, y) = \max(0, 1 + \max_{j \neq y} w^{(j)T} x - w^{(y)T} x)$

7 Kernels

Kernel trick: 1. Express problem s.t. it only depends on inner products $x_j^T x_i$.

2. Replace inner products by kernels.

Reformulate problem: Fundamental insight: Optimal separating hyperplane lives in the span of data: $\hat{w} = \sum_{i=1}^n \alpha_i y_i x_i$

Perceptron example: $\hat{w} = \arg \min_w \frac{1}{n} \sum_{i=1}^n \max(0, -y_i w^T x_i)$

$\hat{a} = \arg \min_{\alpha} \frac{1}{n} \sum_{i=1}^n \max(0, -y_i (\sum_{j=1}^n \alpha_j y_j x_j)^T x_i)$

$\hat{a} = \arg \min_{\alpha} \frac{1}{n} \sum_{i=1}^n \max(0, -\sum_{j=1}^n \alpha_j y_i y_j x_j^T x_i)$

Replace inner products: For some feature transform $\phi: x \mapsto \phi(x)$, kernels solve $\phi(x)^T \phi(x')$ **efficiently** as $k(x, x')$.

Perceptron example (Training):

1. Initialize $\alpha_1 = \alpha_2 = \dots = \alpha_n = 0$

2. For $t = 1, 2, \dots$

(a) Pick $(x_i, y_i) \in D$ u.a.r.

(b) Predict $\hat{y} = \text{sign} \left(\sum_{j=1}^n \alpha_j y_j k(x_j, x_i) \right)$

(c) If $\hat{y} \neq y_i$ set $\alpha_i \leftarrow \alpha_i + \eta_t$

Kernel properties: $k: X \times X \Rightarrow \mathbb{R}$ must be symmetric: $k(x, x') = k(x', x)$

Gram matrix K must be p.s.d. ($\forall x. x^T K x \geq 0$)

For any n , any set $\{x_1, \dots, x_n\} \subseteq X$. All p.s.d. matrices are some kernel and all kernels have Gram matrix.

$$K = \begin{bmatrix} k(x_1, x_1) & \dots & k(x_1, x_n) \\ \vdots & \ddots & \vdots \\ k(x_n, x_1) & \dots & k(x_n, x_n) \end{bmatrix}$$

For k_1, k_2 kernel, $c > 0$ and f polyn. with pos. coef. or exp. $k_1 + k_2, k_1 \cdot k_2, c \cdot k_1$ and $f(k_1(x, x'))$ are kernels.

Poly. degree $= d \quad (x^T x')^d$

Poly. degree $\leq d \quad (x^T x' + 1)^d$

Gaussian (RBF) $\exp(-\|x - x'\|_2^2 / (2h^2))$

Laplacian $\exp(-\|x - x'\|_1 / h)$

Note: $h > 0$ is bandwidth, $h \rightarrow 0$ overfits.

k-NN: $y = \text{sign}(\sum_{i=1}^n y_i k(x_i \text{ among k-NN of } x))$

No training, but **depends on all data**.

$y = \text{sign}(\sum_{i=1}^n y_i \alpha_i k(x_i, x))$

Can use kernel as *similarity function*: **Improved performance**, depends only on **wrongly classified data**, can capture **global trends**, but **requires training**.

Parametric vs. nonparametric learning:

Parametric have finite set of parameters (regression, perceptron), while *nonparametric* increase complexity with size of data (kernelized perceptron, k-NN).

Can kernelize other tasks, such as SVM. Let

$k_i = [y_1 k(x_1, x_i) \quad \dots \quad y_n k(x_n, x_i)]$:

$\arg \min_{\alpha} \frac{1}{n} \sum_{i=1}^n \max\{0, 1 - y_i \alpha^T k_i\} + \lambda \alpha^T D_y K D_y \alpha$

Lin. reg. (training): $\hat{\alpha} =$

$\arg \min_{\alpha} \frac{1}{n} \|\alpha^T K - y\|_2^2 + \lambda \alpha^T K \alpha$

Closed form sol: $\hat{\alpha} = (K + n\lambda I)^{-1} y$

Lin. reg. (prediction): $\hat{y} = \sum_{i=1}^n \hat{\alpha}_i k(x_i, x)$

Est. kernel parameters via CV. Choosing kernels requires **domain knowledge**. Deal w/ overfit by regularization.

8 Neural Networks

Parametrized feature maps + regression. Apply nonlinear activation function φ after weighted

\sum of inputs. Can learn nonlinear features. *Forward propagation* in layer ℓ : ($v^{(0)} = x$): $v^{(\ell)} =$

$\varphi(z^{(\ell)}) = \varphi(W^{(\ell)} v^{(\ell-1)})$

Output layer: $y = f = W^{(L-1)} v^{(L-1)}$. (f can be vector)

Weight optimization: Apply loss $\ell(y - f(x, W))$, optimize weights to minimize loss.

For multi-outputs sum losses.

Sigmoid $\frac{1}{1+\exp(-z)}$

Tanh $\frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$

ReLU $\max(0, z)$

diff. everywhere, not diff. at 0, $\varphi' =$

Back propagation: $\delta^{(L)} = l'(f) =$

$[l'(f_1) \quad \dots \quad l'(f_p)]$, and $\nabla_{W^{(L)}} \ell(W; y, x) =$

$\delta^{(L)} v^{(L-1)T}$. For $\ell < L$: $\delta^{(\ell)} = \varphi'(z^{(\ell)}) \odot$

$(W^{(\ell+1)T} \delta^{(\ell+1)})$ and $\nabla_{w^\ell} \ell(W; y, x) =$

$\delta^{(\ell)} v^{(\ell-1)T}$

Nonconvex optimization, initialization matters!

Random works well.

Glorot (tanh): $w_{i,j} \sim \mathcal{N}(0, 1/n_{in})$

$w_{i,j} \sim \mathcal{N}(0, 2/(n_{in} + n_{out}))$

He (ReLU): $w_{i,j} \sim \mathcal{N}(0, 2/(n_{in}))$

Learning rate η_t decreasing (e.g. $\min(0.1, 100/t)$) to **prevent oscillation**.

Momentum ($m < 1$, $a \leftarrow m \cdot a + \eta_t \nabla_w \ell(W; y, x)$; $W \leftarrow W - a$) can **avoid local minima**.

Many parameters \rightarrow **overfitting!** *Early stop*

or *regularization* ($\lambda \|W\|_F^2$) to prevent. Also

dropout, randomly set weights to 0 with prob. p , set $W = W \odot p$ after training.

Batch normalization, standardize some batch $\{x_{1...m}\}$ and set $y_i = \gamma \hat{x}_i + \beta = BN_{\gamma, \beta}(x_i)$.

Then $\varphi(Wx) = \varphi(W(BN_{\gamma, \beta}(x)))$.

Convolutional NN: Apply $m f \times f$ filters to $n \times n$ image, padding p and stride s : Leaves with $\alpha \times \alpha \times m$ output, where $\alpha = \frac{n+2p-f}{s}$.

9 Unsupervised Learning

k-Means clust.: Represent cluster as center, assign point $x_i \in \mathbb{R}^d$ to nearest center $\mu_j \in \mathbb{R}^d$.

Squared loss (below) **nonconvex**.

$\hat{R}(\mu) = \sum_{i=1}^n \min_{j \in [k]} \|x_i - \mu_j\|_2^2$

Lloyd: Initialize $\mu^{(0)}$. Then assign x_i to closest center $z_i^{(t)} = \arg \min_j \|x_i - \mu_j^{(t-1)}\|_2^2$. Then

update mean: $\mu_j^{(t)} = \frac{1}{|z: z_i^{(t)} = j|} \sum_{i: z_i^{(t)} = j} x_j$

$\mathcal{O}(nkd)$ per iteration, converges **pot. slowly** but **monotonically to local optimum** \rightarrow **Multiple iter.**

k-Means++: Let $\mu^{(0)} = x$ u.a.r. from X . Assign centers $2 \dots k$ randomly, prop. to sq. dist. to closest sel. cent. Expected cost within $\mathcal{O}(\log k)$ of optimum.

$\Pr[\mu_j = x_i] = \frac{1}{k} \min_{k \in [j-1]} \|\mu_k - x_i\|_2^2$

Choosing k **difficult**. Heuristic: When $k + 1$ yields *diminishing returns*, or *regularization* with λk . Only models circular clusters \rightarrow use kernels.

Dimension Reduction: Embed $\{x_1, \dots, x_n\}$, $x_i \in \mathbb{R}^d$ in \mathbb{R}^k where $k < d$.

PCA: Center data $\mu = \frac{1}{n} \sum_{i=1}^n x_i = 0$ and con-

struct **empirical cov. matrix:** $\Sigma = \frac{1}{n} \sum_{i=1}^n x_i x_i^T$

PCA problem: $(W, z_1, \dots, z_n) =$

$\arg \min_{W, z} \sum_{i=1}^n \|W z_i - x_i\|_2^2$ where W or-

thogonal, $z_i \in \mathbb{R}^k$ has sol. $z_i = W^T x_i$.

The solution for W is given by the k principal eigenvectors of Σ . $W = (v_1 | \dots | v_k)$ where

$\Sigma = \sum_{i=1}^d \lambda_i v_i v_i^T$ with $\lambda_1 \geq \dots \geq \lambda_d \geq 0$.

Can apply to any matrix $X = USV^T$, first k components of V are first k principal compo-

nents. Choose k by CV for feature ind., else s.t. variance is explained (see k-Means). Solve

nonlinear PCA via kernels.

Kernel PCA: $w = \sum_{j=1}^n \alpha_j \phi(x_j)$, $K =$

$\sum_{i=1}^n \lambda_i v_i v_i^T$

$\arg \max_{\|w\|_2=1} \sum_{i=1}^n (w^T \phi(x_i))^2 =$

$\arg \max_{\alpha} \alpha^T K \alpha$, Solution: $\alpha^{(i)} =$

$v_i / (\sqrt{\lambda_i})$

New point: $z_i = w^{(i)T} x = \sum_{j=1}^n \alpha_j^{(i)} k(x, x_j)$

Centering a kernel: $K' = K - KE - EK +$

EKE where $E = \frac{1}{n} [1, \dots, 1] [1, \dots, 1]^T$.

Complexity **grows with the number of data points**.

Autoencoder: NN where hidden layers usually smaller (k) than in- and output (d). Try to learn identity function. Compression from input to smallest HL, Decompression from smallest HL to output. $f(x; \theta) = f_{dec}(f_{enc}(x; \theta_1); \theta_2)$. If

$\varphi(z) = z$ then autoencoder is equivalent to PCA.

10 Probabilistic modeling

Bayes optimal predictor: $h^*(x) = \mathbb{E}[Y | X = x]$, **unattainable in pr.** Can try to estimate condi-

tional distr. $\hat{P}(Y | X)$.

Parametric estimation: Have $\hat{P}(Y | X, \theta)$, MLE:

$\theta^* = \arg \max_{\theta} \hat{P}(Y | X, \theta)$

$\arg \min_{\theta} - \sum_{i=1}^n \log \hat{P}(y_i | x_i, \theta)$

Ex.: Gaussian noise, lin. reg.: $y_i \sim$

$\mathcal{N}(w^T x, \sigma^2)$. Then MLE:

$\arg \min_{\theta} \frac{n}{2} \log(2\pi\sigma^2) \sum_{i=1}^n \frac{(y_i - w^T x_i)^2}{2\sigma^2}$

and MLE equivalent to LSQ estimation. In general, MLE with Gaussian noise with constant variance is **equivalent** to LSQ sol.

Bias Variance Tradeoff: Prediction error = Bias² + Variance + Noise.

Bias: Excess risk of best model considered compared to minimal achievable risk knowing

$P(X, Y)$ (i.e., given infinite data).

Variance: Risk incurred due to estimating model from limited data.

Noise: Risk incurred by optimal model (i.e., irreducible error).

Trade bias and variance via model selection /

regularization

Maximum A Posteriori estimation: Placing assumptions on distribution of parameter. For Gaussian noise and Gaussian prior: MAP =

Ridge regression.

Regularized estimation can often be understood as MAP inference:

$\arg \min_w \sum_{i=1}^n \ell(w^T x_i; x_i, y_i) + C(w) =$

$\arg \max_w \prod_i P(y_i | x_i, w) P(w) =$

$\arg \max_w P(w | D)$ where $C(w) = -\log P(w)$

and $\ell(w^T x_i; x_i, y_i) = -\log P(y_i | x_i, w)$.

Bayes optimal classifier: $h^*(x) =$

$\arg \max_y P(Y = y | X = x)$

Logistic regression: $P(Y = y | x) =$

$\frac{1}{1+\exp(-yw^T x)}$. Replaces Gaussian noise as-

sumption with Bernoulli noise: $P(y | x, w) =$

$\text{Ber}(y; \sigma(w^T x))$. MLE: $\hat{R}(w) = \sum_{i=1}^n \log(1 +$

$\exp(-y_i w^T x_i))$, is **convex** and $\nabla_w \ell_w =$

$\frac{1}{1+\exp(-yw^T x)} \exp(-yw^T x)(-yx)$, and

$\exp(-yw^T x) = 1$ if misclassified. Thus

$\nabla_w \ell_w = \frac{-yx}{1+\exp(yw^T x)}$. With L2 regular-

izer, take step in direction $w(1 - 2\lambda\eta_t) -$

$\eta_t \nabla_w \ell_w(y, x)$. Nonlinear classification via

kernels: $\hat{\alpha} = \arg \min_{\alpha} \sum_{i=1}^n \log(1 +$

$\exp(-y_i \alpha^T K_i)) + \lambda \alpha^T K \alpha$ and $\hat{P}(y | x, \hat{\alpha}) =$

$\frac{1}{1+\exp(-y \sum_{j=1}^n \alpha_j k(x_j, x))}$ with $w = \sum_i \alpha_i x_i$.

Multi-class: $P(Y = i | x, w_1, \dots, w_c) =$

$\frac{\exp(w_i^T x)}{\sum_{j=1}^c \exp(w_j^T x)}$

Can obtain class probabilities, but **dense solutions**

11 Decision Theory

Have $P(y | x)$, actions \mathcal{A} and cost $C: \mathcal{Y} \times \mathcal{A} \rightarrow$

\mathbb{R} . Min. Exp. cost: $a^* = \arg \min_a \mathbb{E}_y[C(y, a) |$

$x]$

Asymmetric costs: Assume $A = \{-1, +1\}$,

$c_{FP}, c_{FN} > 0$. Let $\hat{P}(y | x) =$

p , then $\mathbb{E}_y[C(y, +1)] = (1 - p)c_{FP}$ and

$\mathbb{E}_y[C(y, -1)] = pc_{FN}$. Predict +1 when $p >$

$\frac{c_{FP}}{c_{FP} + c_{FN}}$.

Uncertainty sampling: Ask user to label the ex-

ample that we are most uncertain about: $i_t \in$

$\arg \min_i |0.5 - \hat{P}(Y_i | x_i)|$. Active learning **vio-**

lates i.i.d. assumption.

MAP summary: 1. Choose likelihood function

\rightarrow loss function

2. Choose prior \rightarrow regularizer

3. Optimize for MAP parameters, choose hyper-

parameters through cross-validation

4. Make predictions via Bayesian Decision The-

ory

Estimate *joint distribution* $P(X, Y)$ instead of $P(Y | X)$. Cond. distr. can be derived from joint: 1. Estimate $P(Y)$, 2. Estimate $P(x|y)$ for each y , 3. Use Bayes' rule: $P(y|x) = P(y)P(x|y)/P(x)$ where $P(x) = \sum_{y'} P(x, y')$. For $c = 2$ *discriminant function* $f(x) = \log \frac{P(Y=1|x)}{P(Y=-1|x)}$, which is $+1$ if $P(Y = 1 | x) > 0.5$.

Naive Bayes: Model class y as categorical, and features **conditionally independent** given y , e.g. Gaussian NB, assumes $P(x_i|y) = \mathcal{N}(x_i|\mu_{y,i}, \sigma_{y,i}^2)$. Produces linear classifier, equiv. to log. reg. if assumptions are met. $f(x) = w^T x + w_0$ where $w_0 = \log \frac{p_+}{1-p_+} + \sum_{i=1}^d \frac{\mu_{+,i}^2 - \mu_{-,i}^2}{2\sigma_i^2}$ and $w_i = \frac{\mu_{+,i} - \mu_{-,i}}{\sigma_i^2}$. Due to conditional independence assumption, predictions can become **overconfident**. # parameters = $O(cd)$, Complexity (memory + inference) **linear in d**.

Categorical Naive Bayes: $P(X_i = c | Y = y) = \theta_{c|y}^{(i)}$. MLE prior: $\hat{p}_y = \frac{\#y}{n}$, MLE feat. distr.: $\theta_{c|y}^{(i)} = \frac{\text{Count}(X_i=c, Y=y)}{\#y}$. **Requires exponentially (in d) many parameters, Fantastic way to overfit.**

Gaussian Bayes: $P(x|y) = \mathcal{N}(x; \mu_y, \Sigma_y)$. MLE prior: $\hat{p}_y = \frac{\#y}{n}$, MLE feat. distr.: $\hat{\mu}_k = \frac{1}{\#y} \sum_{i:y_i=y} x_i$ and $\hat{\Sigma}_y = \frac{1}{\#y} \sum_{i:y_i=y} (x_i - \hat{\mu}_y)(x_i - \hat{\mu}_y)^T$. Discriminant $f(x)$: $\log \frac{p}{1-p} + \frac{1}{2} \left[\log \frac{|\hat{\Sigma}_-|}{|\hat{\Sigma}_+|} + ((x - \hat{\mu}_-)^T \hat{\Sigma}_-^{-1} (x - \hat{\mu}_-) - (x - \hat{\mu}_+)^T \hat{\Sigma}_+^{-1} (x - \hat{\mu}_+)) \right]$.

Captures correlations among features, avoids overconfidence, # parameters = $O(cd^2)$, complexity quadratic in d.

Fischer's LDA: Assume $p = 0.5$ and $\hat{\Sigma}_- = \hat{\Sigma}_+ = \hat{\Sigma}$. Then $f(x) = x^T \hat{\Sigma}^{-1} (\hat{\mu}_+ - \hat{\mu}_-) + \frac{1}{2} (\hat{\mu}_-^T \hat{\Sigma}^{-1} \hat{\mu}_- - \hat{\mu}_+^T \hat{\Sigma}^{-1} \hat{\mu}_+)$. Produces linear classifier, equiv. to log. reg. if assumptions are met. LDA can be viewed as a projection to a 1-dim. subspace that maximizes ratio of between-class and within-class variances. In constrast, PCA (k=1) maximizes the variance of the resulting 1-dim. projection. Generative model, can be used to **detect outliers**, not very robust against violation of **normality of X assumption**.

Can regularize: Beta($\theta, \alpha_+, \alpha_-$) models likelihood of θ given α_+ weight for $y = 1$ and α_- weight for $y = -1$.

Conjugate distr.: Posterior is same family as

prior. Ex.: Beta($\theta, \alpha_+, \alpha_-$) and Beta($\theta, n_+ + \alpha_+, n_- + \alpha_-$). MAP estimate: $\hat{\theta} = \frac{\alpha_+ + n_+ - 1}{\alpha_+ + n_+ + \alpha_- + n_- - 2}$.

13 Generalized Mixture Models

Setting: Labels potentially unknown, want to cluster data. Model $P(x, \theta)$ as *convex combination* of Gaussian distributions: $\sum_{i=1}^c w_i \mathcal{N}(x; \mu_i, \Sigma_i)$ with $w_i > 0$ and $\sum_i w_i = 1$.

$(\mu^*, \Sigma^*, w^*) = \arg \min - \sum_{i=1}^n \log \sum_{j=1}^k w_j \mathcal{N}(x_i; \mu_j, \Sigma_j)$.

Constraints (Σ p.s.d.) **hard to maintain** in SGD. Fitting a GMM = Training a GBC without labels. **Hard-EM:** E-step: Predict most likely class for all x_i : $z_i^{(t)} = \arg \max_z P(z | x_i, \theta^{(t-1)}) = \arg \max_z P(z | \theta^{(t-1)}) P(x_i | z, \theta^{(t-1)})$ with $\theta^{(t)} = [w_{1:c}^{(t)}, \mu_{1:c}^{(t)}, \Sigma_{1:c}^{(t)}]$. M-step: Compute $\theta^{(t)}$ as MLE as for the Gaussian Bayes classifier: $\theta^{(t)} = \arg \max_{\theta} P(D^{(t)} | \theta)$.

Too much information extracted from each label, **overlapping** clusters not detected, **fixed label** when model uncertain. k-Means Algorithm is equiv to Hard-EM if $w_j = 1/k$ and $\Sigma_j = I \cdot \sigma^2$.

Soft-EM: Let $\gamma_j(x)$ be prob. that x in cluster j : $\gamma_j(x) = P(Z = j | x, \Sigma, \mu, w) = \frac{w_j P(x | \Sigma_j, \mu_j)}{\sum_{\ell} w_{\ell} P(x | \Sigma_{\ell}, \mu_{\ell})}$.

E-Step: Calculate prob $\gamma_j(x_i)$ for all i, j based on $\mu^{(t-1)}, \Sigma^{(t-1)}$ and $w^{(t-1)}$.

M-Step: Adjust parameters: $w_j^{(t)} = \frac{1}{n} \sum_{i=1}^n \gamma_j^{(t)}(x_i)$, $\mu_j^{(t)} = \left(\frac{x - \hat{\mu} \sum_{i=1}^n \gamma_j^{(t)}(x_i)}{\sum_{i=1}^n \gamma_j^{(t)}(x_i)} \right)$ and $\Sigma_j^{(t)} = \frac{\sum_{i=1}^n \gamma_j^{(t)}(x_i) (x_i - \mu_j^{(t)})(x_i - \mu_j^{(t)})^T}{\sum_{i=1}^n \gamma_j^{(t)}(x_i)}$.

Initialization sensitive, nonconvex objective. Init. w with uniform distribution, μ by k-means++ and Σ as spherical (according to empirical variance in data). Choose k **via CV**. Avoid **degeneracy** by **adding term $\nu^2 I$** to $\Sigma_j^{(t)} \rightarrow$ Wishart-prior.

Semi-supervised learning: For points with label y_i : $\gamma_j^{(t)}(x_i) = [j = y_i]$

GANs: Objective: $\min_{w_G} \max_{w_D} \mathbb{E}_{x \sim \text{Data}} \log D(x; w_D) + \mathbb{E}_{x \sim \mathcal{N}} \log(1 - D(G(z; w_G); w_D))$, **Mode Collapse:** generator produces less diverse samples than distribution, **Data memorization**, Simultaneous training \rightarrow **oscillations, Cannot compute likelihood on holdout set.**